

Jiakun Yan

✉ jiakuny3@illinois.edu • 🌐 jiakunyan.github.io

Research Interests

My research interest lies in parallel computing, especially in **high-performance RDMA-based system design**. Currently, I am focusing on designing better communication libraries for highly dynamic/irregular programming systems and applications. I am the main developer of *Lightweight Communication Interface (LCI)* and also contributed/contributing to *MPICH*, *HPX*, *Legion*, and *Charm++*.

Education

University of Illinois at Urbana-Champaign

Illinois, USA

Aug. 2020 - Dec. 2025 (Expected)

- Computer Science PhD student, advised by Marc Snir.
- Research around high-performance communication libraries, especially *Lightweight Communication Interface (LCI)*, investigating their support for asynchronous and multithreaded execution and their implications for emerging programming models.
- GPA: 4.0/4.0

Shanghai Jiao Tong University

Shanghai, China

Sep. 2016 - Jun. 2020

- Bachelor's Degree of Engineering, Dept. of Computer Science.
- Zhiyuan Honors Program of Engineering (an elite program for top 5% talented students)
- GPA: 91.88/100 | Ranking: 4th/151.

University of California, Berkeley

California, USA

Jan. 2019 - May 2019

- Exchange student, Berkeley Global Access Discover Program, GPA: 4.0/4.0.

Experience

GPU Software, Legate Group

NVIDIA

Software Engineer Intern, worked with Manolis Papadakis and Hessam Mirsadeghi May. 2024 - Aug. 2024

- Performance profiling and optimization for *Legion* UCX backend.

Programming Models and Runtime Systems Group

Argonne National Laboratory

Research Intern, advised by Yanfei Guo

May. 2023 - Aug. 2023

- Design and Evaluation of the MPI Continuation and VCI extensions in *MPICH*.

Programming Systems and Applications Research Group

NVIDIA Research

Research Intern, advised by Michael Bauer and Michael Garland

May. 2022 - Aug. 2022

- Realm Collective: design and implement collective communication operations in *Legion* Realm.

PASSION Lab

Lawrence Berkeley Laboratory

Research Assistant, advised by Aydın Buluç and Katherine Yelick

Aug. 2019 - Jan. 2020

- Asynchronous RPC Library (ARL): a high-throughput RPC system with node-level aggregation and single-node work-stealing.
- RDMA vs. RPC for Implementing Distributed Data Structures

Publication

- **Jiakun Yan**, Marc Snir, Yanfei Guo. *Examine MPI and its Extensions for Asynchronous Multithreaded Communication*, Proceedings of the 32nd European MPI Users' Group Meeting (EuroMPI), 2025.
- **Jiakun Yan**, Marc Snir. *LCI: a Lightweight Communication Interface for Efficient Asynchronous Multithreaded Communication*, Proceedings of the 2025 International Conference for High Performance Computing, Networking, Storage and Analysis (SC), 2025.

- **Jiakun Yan**, Hartmut Kaiser, Marc Snir. *Understanding the Communication Needs of Asynchronous Many-Task Systems – A Case Study of HPX+LCI*, preprint, 2025.
- **Jiakun Yan**, Marc Snir. *Contemplating a Lightweight Communication Interface for Asynchronous Many-Task Systems*, Workshop on Asynchronous Many-Task Systems and Applications (WAMTA), 2025.
- Gregor Daiß, Patrick Diehl, **Jiakun Yan**, John K. Holmen, Rahulkumar Gayatri, Christoph Junghans, Alexander Straub, Jeff R. Hammond, Dominic Marcello, Miwako Tsuji, Dirk Pflüger, Hartmut Kaiser. *Asynchronous-Many-Task Systems: Challenges and Opportunities – Scaling an AMR Astrophysics Code on Exascale machines using Kokkos and HPX*, preprint, 2024.
- **Jiakun Yan**, Hartmut Kaiser, and Marc Snir. *Design and Analysis of the Network Software Stack of an Asynchronous Many-task System – The LCI parcellport of HPX*, SC'23 Workshops of The International Conference on High Performance Computing, Network, Storage, and Analysis (SC-W), 2023.
- Benjamin Brock, Yuxin Chen, **Jiakun Yan**, John Owens, Aydın Buluç, and Katherine Yelick. *RDMA vs. RPC for Implementing Distributed Data Structures*, 2019 IEEE/ACM 9th Workshop on Irregular Applications: Architectures and Algorithms (IA³), 2019.

Project

Lightweight Communication Interface

Advised by Marc Snir

UIUC

Aug. 2020 - Present

- *Lightweight Communication Interface (LCI)* is an experimental communication library and research tool designed to support asynchronous communication in multithreaded environments. It features a unified interface for common communication primitives, flexible synchronization semantics, and explicit control over communication resources and runtime behavior.
- LCI delivers optimal multithreaded performance, on par with the traditional one-process-per-core setup. LCI has been integrated into *PaRSEC* and *HPX* and outperforms their original communication backends. We are working on integrating LCI into *Charm++*.

HPX over LCI

Advised by Marc Snir and Hartmut Kaiser

UIUC

Aug. 2021 - Oct. 2023

- *HPX* is a runtime system known for its support for the asynchronous task programming model. In this project, we added an LCI parcellport for *HPX*, enabling more direct support of *HPX* communication and scheduling logic.
- Compared to the MPI parcellport, the LCI parcellport achieves up to 100x improvement in microbenchmark and around 2x in a real-world application, *Octo-Tiger* (on NERSC Perlmutter with 1720 nodes).
- The LCI parcellport has been shipped with *HPX* releases since *HPX* 1.8.0 and used in a few *HPX* applications.

Collective Communication Operations in Realm

Advised by Michael Bauer and Michael Garland

NVIDIA Research

May 2021 - Aug. 2021

- *Realm* is an event-based low-level runtime system providing a high-performance asynchronous task execution model for the higher-level data-centric parallel programming system *Legion*. In this project, we extended its generic copy operation to handle collective broadcast communication across CPU/GPU buffers.
- We designed and implemented a hierarchical path planning algorithm that includes inter-node radix tree broadcast and intra-node path aggregation. It achieved significant improvement compared to the original point-to-point copies on a set of synthetic benchmarks.

Honors and Awards

- **Best Poster Award**, WAMTA24 Feb. 2024
A Lightweight Communication Interface for Asynchronous Many-Task Systems

Skills

- **Programming Language:** C, C++, Python, CUDA, Java, Rust, Go
- **Library & Framework:** libibverbs, libfabric, UCX, MPI, GASNet-EX, UPC++, OpenSHMEM, Argobots